

In [200]:

```
# here we are importing the all necessary packages
from imblearn.under_sampling import RandomUnderSampler

import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import plotly.offline as py
py.init_notebook_mode(connected=True)
import plotly.graph_objs as go
import plotly.tools as tls
import os
import gc
import pandas as pd
import matplotlib.pyplot as plt
import time
import warnings
import numpy as np
from sklearn.ensemble import RandomForestClassifier
warnings.filterwarnings("ignore")
from sklearn import model_selection
from sklearn.linear_model import LogisticRegression
from scipy import stats
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from imblearn.over_sampling import SMOTE
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from imblearn.pipeline import Pipeline
from sklearn.model_selection import train_test_split
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import confusion_matrix
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
import lightgbm as lgb
from sklearn.svm import SVC
# using one hot encoding
from sklearn.preprocessing import OneHotEncoder
from sklearn import metrics
from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2
from sklearn import metrics
import tensorflow
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.models import Model
```

In [201]:

```
# code for printing the confusion matrix
#https://gist.github.com/shaypal5/94c53d765083101efc0240d776a23823i
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

def print_confusion_matrix(confusion_matrix, class_names, figsize = (10,7), fontsize=14):

    df_cm = pd.DataFrame(
        confusion_matrix, index=class_names, columns=class_names,
    )
    fig = plt.figure(figsize=figsize)
    try:
        heatmap = sns.heatmap(df_cm, annot=True, fmt="d")
    except ValueError:
        raise ValueError("Confusion matrix values must be integers.")
    heatmap.yaxis.set_ticklabels(heatmap.yaxis.get_ticklabels(), rotation=0, ha='right', fo
    heatmap.xaxis.set_ticklabels(heatmap.xaxis.get_ticklabels(), rotation=0, ha='right', fo
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
```

In [202]:

```
# reading the data using pandas
data=pd.read_csv('/content/drive/My Drive/Artificial engineering/diabetic_data.csv')
print("Number of data points:",data.shape[0])

condition = data['readmitted']=='<30'
data['readmitted'] = np.where(condition,1,0)
```

Number of data points: 101766

In [203]:

```
# splitting into x and y
y=data['readmitted']
X=data.drop(columns=['readmitted'])
```

In [204]:

```
# printing the data columns
data.columns
```

Out[204]:

```
Index(['encounter_id', 'patient_nbr', 'race', 'gender', 'age', 'weight',
      'admission_type_id', 'discharge_disposition_id', 'admission_source_id',
      'time_in_hospital', 'payer_code', 'medical_specialty',
      'num_lab_procedures', 'num_procedures', 'num_medications',
      'number_outpatient', 'number_emergency', 'number_inpatient', 'diag_1',
      'diag_2', 'diag_3', 'number_diagnoses', 'max_glu_serum', 'A1Cresult',
      'metformin', 'repaglinide', 'nateglinide', 'chlorpropamide',
      'glimepiride', 'acetoexamide', 'glipizide', 'glyburide', 'tolbutamide',
      'pioglitazone', 'rosiglitazone', 'acarbose', 'miglitol', 'troglitazone',
      'tolazamide', 'examide', 'citoglipton', 'insulin',
      'glyburide-metformin', 'glipizide-metformin',
      'glimepiride-pioglitazone', 'metformin-rosiglitazone',
      'metformin-pioglitazone', 'change', 'diabetesMed', 'readmitted'],
      dtype='object')
```

In [205]:

```
# printing the data types  
data.dtypes
```

Out[205]:

```
encounter_id          int64  
patient_nbr           int64  
race                  object  
gender                object  
age                  object  
weight                object  
admission_type_id     int64  
discharge_disposition_id int64  
admission_source_id   int64  
time_in_hospital      int64  
payer_code            object  
medical_specialty      object  
num_lab_procedures    int64  
num_procedures         int64  
num_medications        int64  
number_outpatient      int64  
number_emergency       int64  
number_inpatient       int64  
diag_1                object  
diag_2                object  
diag_3                object  
number_diagnoses       int64  
max_glu_serum          object  
A1Cresult              object  
metformin              object  
repaglinide            object  
nateglinide            object  
chlorpropamide         object  
glimepiride            object  
acetohexamide          object  
glipizide              object  
glyburide              object  
tolbutamide            object  
pioglitazone           object  
rosiglitazone          object  
acarbose               object  
miglitol               object  
troglitazone           object  
tolazamide             object  
examide                object  
citoglipton            object  
insulin                object  
glyburide-metformin    object  
glipizide-metformin    object  
glimepiride-pioglitazone object  
metformin-rosiglitazone object  
metformin-pioglitazone object  
change                 object  
diabetesMed            object  
readmitted             int64  
dtype: object
```

In [206]:

```
# splitting  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, stratify=y)
```

In [207]:

```
# taking all categorical data and convert into one hot encoding
enc = OneHotEncoder(handle_unknown='ignore')
enc.fit(X_train['race'].values.reshape(-1,1))
X_train_race=enc.transform(X_train['race'].values.reshape(-1,1))
X_test_race=enc.transform(X_test['race'].values.reshape(-1,1))
enc.fit(X_train['gender'].values.reshape(-1,1))
X_train_gender=enc.transform(X_train['gender'].values.reshape(-1,1))
X_test_gender=enc.transform(X_test['gender'].values.reshape(-1,1))
enc.fit(X_train['age'].values.reshape(-1,1))
X_train_age=enc.transform(X_train['age'].values.reshape(-1,1))
X_test_age=enc.transform(X_test['age'].values.reshape(-1,1))
enc.fit(X_train['weight'].values.reshape(-1,1))
X_train_weight=enc.transform(X_train['weight'].values.reshape(-1,1))
X_test_weight=enc.transform(X_test['weight'].values.reshape(-1,1))
enc.fit(X_train['payer_code'].values.reshape(-1,1))
X_train_payer_code=enc.transform(X_train['payer_code'].values.reshape(-1,1))
X_test_payer_code=enc.transform(X_test['payer_code'].values.reshape(-1,1))
enc.fit(X_train['medical_specialty'].values.reshape(-1,1))
X_train_medical_specialty=enc.transform(X_train['medical_specialty'].values.reshape(-1,1))
X_test_medical_specialty=enc.transform(X_test['medical_specialty'].values.reshape(-1,1))
enc.fit(X_train['diag_1'].values.reshape(-1,1))
X_train_diag_1=enc.transform(X_train['diag_1'].values.reshape(-1,1))
X_test_diag_1=enc.transform(X_test['diag_1'].values.reshape(-1,1))
enc.fit(X_train['diag_2'].values.reshape(-1,1))
X_train_diag_2=enc.transform(X_train['diag_2'].values.reshape(-1,1))
X_test_diag_2=enc.transform(X_test['diag_2'].values.reshape(-1,1))
enc.fit(X_train['diag_3'].values.reshape(-1,1))
X_train_diag_3=enc.transform(X_train['diag_3'].values.reshape(-1,1))
X_test_diag_3=enc.transform(X_test['diag_3'].values.reshape(-1,1))
enc.fit(X_train['max_glu_serum'].values.reshape(-1,1))
X_train_max_glu_serum=enc.transform(X_train['max_glu_serum'].values.reshape(-1,1))
X_test_max_glu_serum=enc.transform(X_test['max_glu_serum'].values.reshape(-1,1))
enc.fit(X_train['A1Cresult'].values.reshape(-1,1))
X_train_A1Cresult=enc.transform(X_train['A1Cresult'].values.reshape(-1,1))
X_test_A1Cresult=enc.transform(X_test['A1Cresult'].values.reshape(-1,1))
enc.fit(X_train['metformin'].values.reshape(-1,1))
X_train_metformin=enc.transform(X_train['metformin'].values.reshape(-1,1))
X_test_metformin=enc.transform(X_test['metformin'].values.reshape(-1,1))
enc.fit(X_train['metformin'].values.reshape(-1,1))
X_train_metformin=enc.transform(X_train['metformin'].values.reshape(-1,1))
X_test_metformin=enc.transform(X_test['metformin'].values.reshape(-1,1))
enc.fit(X_train['repaglinide'].values.reshape(-1,1))
X_train_repaglinide=enc.transform(X_train['repaglinide'].values.reshape(-1,1))
X_test_repaglinide=enc.transform(X_test['repaglinide'].values.reshape(-1,1))
enc.fit(X_train['nateglinide'].values.reshape(-1,1))
X_train_nateglinide=enc.transform(X_train['nateglinide'].values.reshape(-1,1))
X_test_nateglinide=enc.transform(X_test['nateglinide'].values.reshape(-1,1))
enc.fit(X_train['chlorpropamide'].values.reshape(-1,1))
X_train_chlorpropamide=enc.transform(X_train['chlorpropamide'].values.reshape(-1,1))
X_test_chlorpropamide=enc.transform(X_test['chlorpropamide'].values.reshape(-1,1))
enc.fit(X_train['glimepiride'].values.reshape(-1,1))
X_train_glimepiride=enc.transform(X_train['glimepiride'].values.reshape(-1,1))
X_test_glimepiride=enc.transform(X_test['glimepiride'].values.reshape(-1,1))
enc.fit(X_train['acetohexamide'].values.reshape(-1,1))
X_train_acetohexamide=enc.transform(X_train['acetohexamide'].values.reshape(-1,1))
X_test_acetohexamide=enc.transform(X_test['acetohexamide'].values.reshape(-1,1))
enc.fit(X_train['glipizide'].values.reshape(-1,1))
X_train_glipizide=enc.transform(X_train['glipizide'].values.reshape(-1,1))
X_test_glipizide=enc.transform(X_test['glipizide'].values.reshape(-1,1))
```

```

enc.fit(X_train['glyburide'].values.reshape(-1,1))
X_train_glyburide=enc.transform(X_train['glyburide'].values.reshape(-1,1))
X_test_glyburide=enc.transform(X_test['glyburide'].values.reshape(-1,1))
enc.fit(X_train['tolbutamide'].values.reshape(-1,1))
X_train_tolbutamide=enc.transform(X_train['tolbutamide'].values.reshape(-1,1))
X_test_tolbutamide=enc.transform(X_test['tolbutamide'].values.reshape(-1,1))
enc.fit(X_train['pioglitazone'].values.reshape(-1,1))
X_train_pioglitazone=enc.transform(X_train['pioglitazone'].values.reshape(-1,1))
X_test_pioglitazone=enc.transform(X_test['pioglitazone'].values.reshape(-1,1))
enc.fit(X_train['rosiglitazone'].values.reshape(-1,1))
X_train_rosiglitazone=enc.transform(X_train['rosiglitazone'].values.reshape(-1,1))
X_test_rosiglitazone=enc.transform(X_test['rosiglitazone'].values.reshape(-1,1))
enc.fit(X_train['acarbose'].values.reshape(-1,1))
X_train_acarbose=enc.transform(X_train['acarbose'].values.reshape(-1,1))
X_test_acarbose=enc.transform(X_test['acarbose'].values.reshape(-1,1))
enc.fit(X_train['miglitol'].values.reshape(-1,1))
X_train_miglitol=enc.transform(X_train['miglitol'].values.reshape(-1,1))
X_test_miglitol=enc.transform(X_test['miglitol'].values.reshape(-1,1))
enc.fit(X_train['troglitazone'].values.reshape(-1,1))
X_train_troglitazone=enc.transform(X_train['troglitazone'].values.reshape(-1,1))
X_test_troglitazone=enc.transform(X_test['troglitazone'].values.reshape(-1,1))
enc.fit(X_train['tolazamide'].values.reshape(-1,1))
X_train_tolazamide=enc.transform(X_train['tolazamide'].values.reshape(-1,1))
X_test_tolazamide=enc.transform(X_test['tolazamide'].values.reshape(-1,1))
enc.fit(X_train['examide'].values.reshape(-1,1))
X_train_examide=enc.transform(X_train['examide'].values.reshape(-1,1))
X_test_examide=enc.transform(X_test['examide'].values.reshape(-1,1))
enc.fit(X_train['citoglipton'].values.reshape(-1,1))
X_train_citoglipton=enc.transform(X_train['citoglipton'].values.reshape(-1,1))
X_test_citoglipton=enc.transform(X_test['citoglipton'].values.reshape(-1,1))
enc.fit(X_train['insulin'].values.reshape(-1,1))
X_train_insulin=enc.transform(X_train['insulin'].values.reshape(-1,1))
X_test_insulin=enc.transform(X_test['insulin'].values.reshape(-1,1))
enc.fit(X_train['glyburide-metformin'].values.reshape(-1,1))
X_train_glyburide_metformin=enc.transform(X_train['glyburide-metformin'].values.reshape(-1,1))
X_test_glyburide_metformin=enc.transform(X_test['glyburide-metformin'].values.reshape(-1,1))
enc.fit(X_train['glipizide-metformin'].values.reshape(-1,1))
X_train_glipizide_metformin=enc.transform(X_train['glipizide-metformin'].values.reshape(-1,1))
X_test_glipizide_metformin=enc.transform(X_test['glipizide-metformin'].values.reshape(-1,1))
enc.fit(X_train['metformin-rosiglitazone'].values.reshape(-1,1))
X_train_metformin_rosiglitazone=enc.transform(X_train['metformin-rosiglitazone'].values.reshape(-1,1))
X_test_metformin_rosiglitazone=enc.transform(X_test['metformin-rosiglitazone'].values.reshape(-1,1))
enc.fit(X_train['glimepiride-pioglitazone'].values.reshape(-1,1))
X_train_glimepiride_pioglitazone=enc.transform(X_train['glimepiride-pioglitazone'].values.reshape(-1,1))
X_test_glimepiride_pioglitazone=enc.transform(X_test['glimepiride-pioglitazone'].values.reshape(-1,1))
enc.fit(X_train['metformin-rosiglitazone'].values.reshape(-1,1))
X_train_metformin_rosiglitazone=enc.transform(X_train['metformin-rosiglitazone'].values.reshape(-1,1))
X_test_metformin_rosiglitazone=enc.transform(X_test['metformin-rosiglitazone'].values.reshape(-1,1))
enc.fit(X_train['metformin-pioglitazone'].values.reshape(-1,1))
X_train_metformin_pioglitazone=enc.transform(X_train['metformin-pioglitazone'].values.reshape(-1,1))
X_test_metformin_pioglitazone=enc.transform(X_test['metformin-pioglitazone'].values.reshape(-1,1))
enc.fit(X_train['change'].values.reshape(-1,1))
X_train_change=enc.transform(X_train['change'].values.reshape(-1,1))
X_test_change=enc.transform(X_test['change'].values.reshape(-1,1))
enc.fit(X_train['diabetesMed'].values.reshape(-1,1))
X_train_diabetesMed=enc.transform(X_train['diabetesMed'].values.reshape(-1,1))
X_test_diabetesMed=enc.transform(X_test['diabetesMed'].values.reshape(-1,1))
enc.fit(X_train['admission_type_id'].values.reshape(-1,1))
X_train_admission_type_id=enc.transform(X_train['admission_type_id'].values.reshape(-1,1))
X_test_admission_type_id=enc.transform(X_test['admission_type_id'].values.reshape(-1,1))
enc.fit(X_train['discharge_disposition_id'].values.reshape(-1,1))

```



```
X_train_discharge_disposition_id=enc.transform(X_train['discharge_disposition_id'].values.reshape(-1,1))
X_test_discharge_disposition_id=enc.transform(X_test['discharge_disposition_id'].values.reshape(-1,1))
enc.fit(X_train['admission_source_id'].values.reshape(-1,1))
X_train_admission_source_id=enc.transform(X_train['admission_source_id'].values.reshape(-1,1))
X_test_admission_source_id=enc.transform(X_test['admission_source_id'].values.reshape(-1,1))
```

In [208]:

```
# normalising the numerical data
from sklearn.preprocessing import Normalizer
normalizer = Normalizer()
normalizer.fit(X_train['patient_nbr'].values.reshape(1,-1))
X_train_patient_nbr = normalizer.transform(X_train['patient_nbr'].values.reshape(1,-1))
X_test_patient_nbr = normalizer.transform(X_test['patient_nbr'].values.reshape(1,-1))
normalizer.fit(X_train['time_in_hospital'].values.reshape(1,-1))
X_train_time_in_hospital = normalizer.transform(X_train['time_in_hospital'].values.reshape(1,-1))
X_test_time_in_hospital = normalizer.transform(X_test['time_in_hospital'].values.reshape(1,-1))
normalizer.fit(X_train['num_lab_procedures'].values.reshape(1,-1))
X_train_num_lab_procedures = normalizer.transform(X_train['num_lab_procedures'].values.reshape(1,-1))
X_test_num_lab_procedures = normalizer.transform(X_test['num_lab_procedures'].values.reshape(1,-1))
normalizer.fit(X_train['num_procedures'].values.reshape(1,-1))
X_train_num_procedures = normalizer.transform(X_train['num_procedures'].values.reshape(1,-1))
X_test_num_procedures = normalizer.transform(X_test['num_procedures'].values.reshape(1,-1))
normalizer.fit(X_train['num_medications'].values.reshape(1,-1))
X_train_num_medications = normalizer.transform(X_train['num_medications'].values.reshape(1,-1))
X_test_num_medications = normalizer.transform(X_test['num_medications'].values.reshape(1,-1))
normalizer.fit(X_train['number_outpatient'].values.reshape(1,-1))
X_train_number_outpatient = normalizer.transform(X_train['number_outpatient'].values.reshape(1,-1))
X_test_number_outpatient = normalizer.transform(X_test['number_outpatient'].values.reshape(1,-1))
normalizer.fit(X_train['number_emergency'].values.reshape(1,-1))
X_train_number_emergency = normalizer.transform(X_train['number_emergency'].values.reshape(1,-1))
X_test_number_emergency = normalizer.transform(X_test['number_emergency'].values.reshape(1,-1))
normalizer.fit(X_train['number_inpatient'].values.reshape(1,-1))
X_train_number_inpatient = normalizer.transform(X_train['number_inpatient'].values.reshape(1,-1))
X_test_number_inpatient = normalizer.transform(X_test['number_inpatient'].values.reshape(1,-1))
normalizer.fit(X_train['number_diagnoses'].values.reshape(1,-1))
X_train_number_diagnoses = normalizer.transform(X_train['number_diagnoses'].values.reshape(1,-1))
X_test_number_diagnoses = normalizer.transform(X_test['number_diagnoses'].values.reshape(1,-1))
X_train_patient_nbr = X_train_patient_nbr.reshape(-1,1)
X_test_patient_nbr = X_test_patient_nbr.reshape(-1,1)
X_train_time_in_hospital = X_train_time_in_hospital.reshape(-1,1)
X_test_time_in_hospital = X_test_time_in_hospital.reshape(-1,1)
X_train_num_lab_procedures = X_train_num_lab_procedures.reshape(-1,1)
X_test_num_lab_procedures = X_test_num_lab_procedures.reshape(-1,1)
X_train_num_procedures = X_train_num_procedures.reshape(-1,1)
X_test_num_procedures = X_test_num_procedures.reshape(-1,1)
X_train_num_medications = X_train_num_medications.reshape(-1,1)
X_test_num_medications = X_test_num_medications.reshape(-1,1)
X_train_number_outpatient = X_train_number_outpatient.reshape(-1,1)
X_test_number_outpatient = X_test_number_outpatient.reshape(-1,1)
X_train_number_emergency = X_train_number_emergency.reshape(-1,1)
X_test_number_emergency = X_test_number_emergency.reshape(-1,1)
X_train_number_inpatient = X_train_number_inpatient.reshape(-1,1)
X_test_number_inpatient = X_test_number_inpatient.reshape(-1,1)
X_train_number_diagnoses = X_train_number_diagnoses.reshape(-1,1)
X_test_number_diagnoses = X_test_number_diagnoses.reshape(-1,1)
```



In [209]:

```
# stacking
from scipy.sparse import hstack
X_train = hstack((X_train_race,X_train_gender,X_train_age,X_train_weight,X_train_payer_code
X_test = hstack((X_test_race,X_test_gender,X_test_age,X_test_weight,X_test_payer_code,X_test
```

In [210]:

```
#https://stackoverflow.com/a/45386397/13693423
# function to select k best features which returns new dataset
def select_features(X_train, y_train, X_test):
    fs = SelectKBest(score_func=chi2, k=150)
    fs.fit(X_train, y_train)
    X_train_fs = fs.transform(X_train)
    X_test_fs = fs.transform(X_test)
    return X_train_fs, X_test_fs, fs
```

In [211]:

```
# to get the best features
from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2
X_train_fs, X_test_fs, fs = select_features(X_train, y_train, X_test)
```

In [212]:

```
# selecting important features
X_train_fs=X_train_fs.toarray()
X_test_fs=X_test_fs.toarray()
```

## Logistic regression

In [213]:

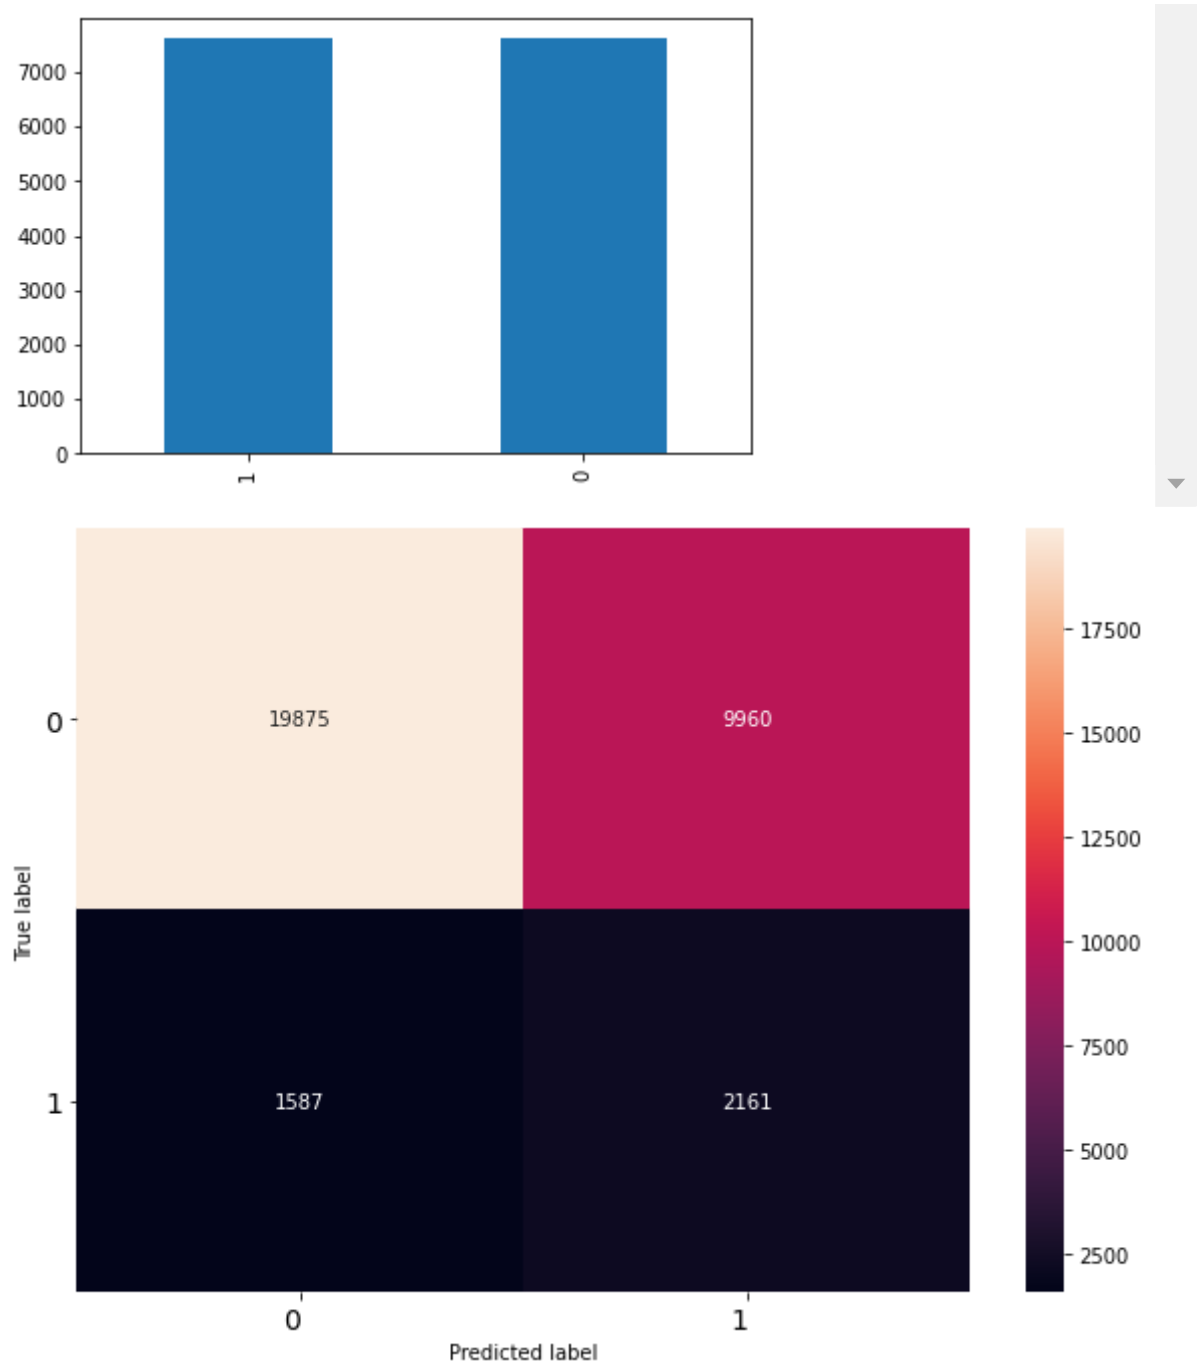
```

from sklearn.model_selection import train_test_split
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import confusion_matrix
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
# function to try on different sets of ratio for smote

# here upsampling is done for imbalanced data
under = RandomUnderSampler()
X_train_fs, y_train = under.fit_resample(X_train_fs, y_train.ravel())
#here bar plot shows equal distribution after upsampling
pd.Series(y_train).value_counts().plot.bar()
# here the scaling is done x train
scaler = StandardScaler()
X_train_fs=scaler.fit_transform(X_train_fs)
X_test_fs=scaler.transform(X_test_fs)
# here using cross validation and gridsearch to determine the best hyperparameter
logit = LogisticRegression()
param_grid = {'C': [0.001, 0.01, 0.1, 1, 10, 100, 1000] }
clf = GridSearchCV(logit, param_grid,scoring='roc_auc',cv=5,n_jobs=-1)
clf.fit(X_train_fs,y_train)
# selecting the best model and printing the confusion matrix
logit =clf.best_estimator_
logit_pred = logit.predict(X_test_fs)
logit_pred_train=logit.predict(X_train_fs)
print_confusion_matrix(confusion_matrix(y_test,logit_pred),['0','1'])
print("Logistic regression")
print("Test confusion matrix")
# printing the accuracy preicison,recall,f1-score from Logistic regression
fpr, tpr, thresholds = metrics.roc_curve(y_test,logit_pred)
print('auc is {0:.2f}'.format(metrics.auc(fpr, tpr)))
log_auc=format(metrics.auc(fpr, tpr))
log_f1_score=format(f1_score(y_test, logit_pred))
print("Accuracy is {0:.2f}".format(accuracy_score(y_test, logit_pred)))
print("Precision is {0:.2f}".format(precision_score(y_test, logit_pred)))
print("Recall is {0:.2f}".format(recall_score(y_test, logit_pred)))
print("f1-score is {0:.2f}".format(f1_score(y_test, logit_pred)))

```

Logistic regression  
 Test confusion matrix  
 auc is 0.62  
 Accuracy is 0.66  
 Precision is 0.18  
 Recall is 0.58  
 f1-score is 0.27



# Decision tree

In [214]:

```


# decision tree
# here using cross validation and gridsearch to determine the best hyperparameter
parameters = {'max_depth': [1, 5, 10, 50], 'min_samples_split': [5, 10, 100, 500]}
dtree = DecisionTreeClassifier()
clf = GridSearchCV(dtree, parameters, scoring='roc_auc', cv=5, n_jobs=-1)
clf.fit(X_train_fs, y_train)
# selecting the best model and printing the confusion matrix
dtree = clf.best_estimator_
dtree.fit(X_train_fs, y_train)
dtree_pred = dtree.predict(X_test_fs)

print_confusion_matrix(confusion_matrix(y_test, dtree_pred), ['0', '1'])
print("Decision tree")
print("Test confusion matrix")

# printing the accuracy preicison, recall, f1-score from Logistic regression
fpr, tpr, thresholds = metrics.roc_curve(y_test, dtree_pred)
print('auc is {0:.2f}'.format(metrics.auc(fpr, tpr)))
decision_auc = format(metrics.auc(fpr, tpr))
decision_f1 = format(f1_score(y_test, dtree_pred))
print("Accuracy is {0:.2f}".format(accuracy_score(y_test, dtree_pred)))
print("Precision is {0:.2f}".format(precision_score(y_test, dtree_pred)))
print("Recall is {0:.2f}".format(recall_score(y_test, dtree_pred)))
print("f1 is {0:.2f}".format(f1_score(y_test, dtree_pred)))

```

Decision tree  
 Test confusion matrix  
 auc is 0.62  
 Accuracy is 0.64  
 Precision is 0.17  
 Recall is 0.58  
 f1 is 0.27



# Random forest

In [215]:

```

# doing hyperparameter tuning for random forest
rm = RandomForestClassifier()
params={'n_estimators':[5,10,25,50,100,300,500], 'n_estimators': [10, 25], 'max_features': [
    'max_depth': [10, 50,75,100, None], 'bootstrap': [True, False]}
model_rf=GridSearchCV(rm,param_grid=params,cv=5,scoring='roc_auc',n_jobs=-1,verbose=1)
model_rf.fit(X_train_fs, y_train)
rm =model_rf.best_estimator_
rm.fit(X_train_fs,y_train)
rm_prd = rm.predict(X_test_fs)

# selecting the best model and printing the confusion matrix
print_confusion_matrix(confusion_matrix(y_test,rm_prd),['0','1'])
print("random forest")
print("Test confusion matrix")
# printing the accuracy preicison,recall,f1-score from Logistic regression
fpr, tpr, thresholds = metrics.roc_curve(y_test,rm_prd)
print('auc is {0:.2f}'.format(metrics.auc(fpr, tpr)))
print("Accuracy is {0:.2f}".format(accuracy_score(y_test, rm_prd)))
print("Precision is {0:.2f}".format(precision_score(y_test, rm_prd)))
print("Recall is {0:.2f}".format(recall_score(y_test, rm_prd)))
print("f1score is {0:.2f}".format(f1_score(y_test, rm_prd)))
rm_auc=format(metrics.auc(fpr, tpr))
rm_f1=format(f1_score(y_test, rm_prd))

```

Fitting 5 folds for each of 40 candidates, totalling 200 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done 64 tasks      | elapsed: 19.0s
[Parallel(n_jobs=-1)]: Done 200 out of 200 | elapsed: 1.4min finished


```

```

random forest
Test confusion matrix
auc is 0.62
Accuracy is 0.63
Precision is 0.17
Recall is 0.61
f1score is 0.27

```






## feature importance

In [216]:

```
# feature importance plot
from matplotlib import pyplot
#for i in range(len(fs.scores_)):
#    print('Feature %d: %f' % (i, fs.scores_[i]))
# plot the scores
pyplot.bar([i for i in range(len(fs.scores_))], fs.scores_)
pyplot.show()
```



## predicting the probability of the model using the decision tree and random forest

In [217]:

```

# in the believe function i wrote the code for weighted average with markdown
from sklearn.model_selection import train_test_split
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import confusion_matrix
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
import lightgbm as lgb
from sklearn.calibration import CalibratedClassifierCV
# function to do weighted average
# used decision tree
decision = DecisionTreeClassifier()
param_grid = {'max_depth': [1, 5, 10, 50], 'min_samples_split': [5, 10, 100, 500]}
clf = GridSearchCV(decision, param_grid, scoring='roc_auc', cv=5, n_jobs=-1)
clf.fit(X_train_fs, y_train)
decision = clf.best_estimator_
sig_clf = CalibratedClassifierCV(decision, method="sigmoid")
sig_clf.fit(X_train_fs, y_train)
# used calibrated classifier to get the probability using decision tree
predict_y_decision = sig_clf.predict_proba(X_test_fs)
# used random forest classifier
rm = RandomForestClassifier()
params={'n_estimators': [5, 10, 25, 50, 100, 300, 500], 'n_estimators': [10, 25], 'max_features': [
    'max_depth': [10, 50, 75, 100, None], 'bootstrap': [True, False]}
model_rf = GridSearchCV(rm, param_grid=params, cv=5, scoring='roc_auc', n_jobs=-1, verbose=1)
model_rf.fit(X_train_fs, y_train)
rm = model_rf.best_estimator_
sig_clf = CalibratedClassifierCV(rm, method="sigmoid")
sig_clf.fit(X_train_fs, y_train)
# used calibrated classifier to get the correct probability predictions
predict_y_rm = sig_clf.predict_proba(X_test_fs)
one, two = predict_y_decision, predict_y_rm
z = []
# added the probability of both the predictions
for i in range(len(one)):
    z.append([one[i][0] + two[i][0], one[i][1] + two[i][1]])
predicted = []
# using argmax got the output and checked for the answer
for i in range(len(z)):
    predicted.append(np.argmax(z[i]))
print_confusion_matrix(confusion_matrix(y_test, predicted), ['0', '1'])
print("combined accuracy")
print("Test confusion matrix")
# printing the accuracy precision, recall, f1-score from Logistic regression
fpr, tpr, thresholds = metrics.roc_curve(y_test, predicted)
print('auc is {0:.2f}'.format(metrics.auc(fpr, tpr)))
print("Precision is {0:.2f}".format(precision_score(y_test, predicted)))
print("Recall is {0:.2f}".format(recall_score(y_test, predicted)))
print("f1score is {0:.2f}".format(f1_score(y_test, predicted)))
avgweight_auc = format(metrics.auc(fpr, tpr))
avgweight_f1 = format(f1_score(y_test, predicted))

```

Fitting 5 folds for each of 40 candidates, totalling 200 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done 64 tasks      | elapsed: 17.8s

```

```
[Parallel(n_jobs=-1)]: Done 200 out of 200 | elapsed: 1.3min finished
```

combined accuracy


Test confusion matrix

auc is 0.63

Precision is 0.18

Recall is 0.60

f1score is 0.27



In [219]:

```

from prettytable import PrettyTable
#If you get a ModuleNotFoundError error , install prettytable using: pip3 install prettytable
x = PrettyTable()
x.field_names = ["Vectorizer", "Model", "f1-score", "TEST AUC"]
x.add_row(["ONE HOT ENCODING", "LOGISTIC REGRESSION", round(float(log_f1_score), 2), round(float(decision_f1), 2)])
x.add_row(["ONE HOT ENCODING", "DECISION TREE", round(float(decision_f1), 2), round(float(decision_f1), 2)])
x.add_row(["ONE HOT ENCODING", "RANDOM FOREST", round(float(rm_f1), 2), round(float(rm_auc), 2)])
x.add_row(["ONE HOT ENCODING", "COMBINED DECISION TREE RANDOM FOREST", round(float(avgweight_f1), 2), round(float(avgweight_auc), 2)])
print(x)

```

```

+-----+-----+-----+-----+
| Vectorizer | Model | f1-score | TEST AUC |
+-----+-----+-----+-----+
| ONE HOT ENCODING | LOGISTIC REGRESSION | 0.27 | 0.6 |
| ONE HOT ENCODING | DECISION TREE | 0.27 | 0.6 |
| ONE HOT ENCODING | RANDOM FOREST | 0.27 | 0.6 |
| ONE HOT ENCODING | COMBINED DECISION TREE RANDOM FOREST | 0.27 | 0.6 |
+-----+-----+-----+-----+

```

## predicting using deep learning

In [ ]:

```

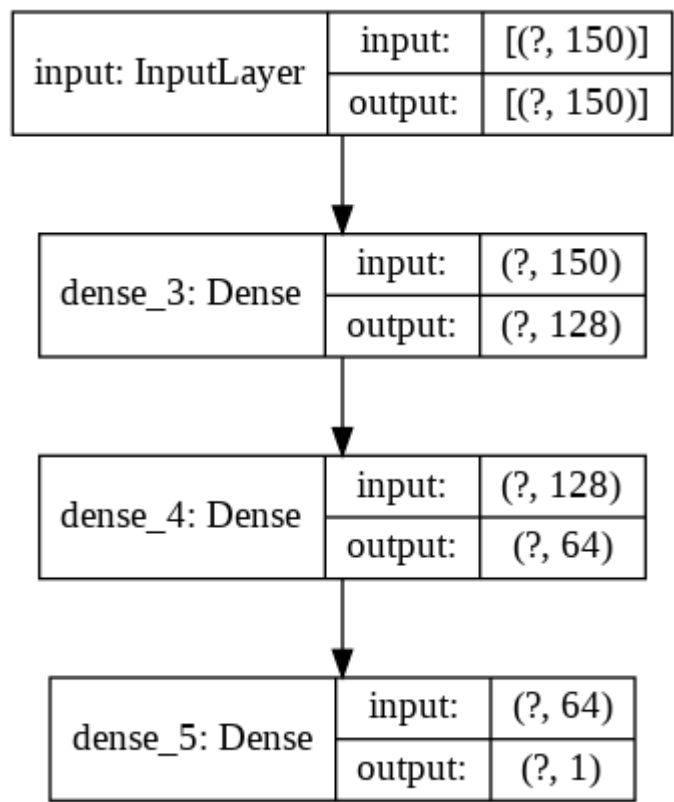
# simple deep learning model
k = tensorflow.keras.initializers.he_uniform(seed=None)
inputlayer=Input(shape=(X_test_fs.shape[1]),name='input')
dense=Dense(128,activation='relu',kernel_initializer=k)(inputlayer)
dense1=Dense(64,activation='relu',kernel_initializer=k)(dense)
output=Dense(1, activation='sigmoid',kernel_initializer=k)(dense1)
model=Model(inputs=inputlayer,outputs=output)
opt = tensorflow.keras.optimizers.Adam(learning_rate=0.001)
model.compile(loss='binary_crossentropy', optimizer=opt, metrics=['accuracy',tensorflow.keras.metrics.AUC()])

```

In [ ]:

```
# plt of the model simple dl model
from tensorflow.keras.utils import plot_model
plot_model(model, to_file='model_plot.png', show_shapes=True, show_layer_names=True)
```

Out[210]:



In [ ]:

```
model= model.fit(X_train_fs,y_train,validation_data=(X_test_fs, y_test), epochs=5)
```

Epoch 1/5

476/476 [=====] - 6s 12ms/step - loss: 0.6718 - accuracy: 0.5989 - auc\_1: 0.6413 - val\_loss: 0.7363 - val\_accuracy: 0.5803 - val\_auc\_1: 0.6565

Epoch 2/5

476/476 [=====] - 5s 11ms/step - loss: 0.6259 - accuracy: 0.6452 - auc\_1: 0.7015 - val\_loss: 0.6711 - val\_accuracy: 0.6526 - val\_auc\_1: 0.6557

Epoch 3/5

476/476 [=====] - 5s 11ms/step - loss: 0.6116 - accuracy: 0.6595 - auc\_1: 0.7203 - val\_loss: 0.6524 - val\_accuracy: 0.6608 - val\_auc\_1: 0.6534

Epoch 4/5

476/476 [=====] - 5s 11ms/step - loss: 0.5998 - accuracy: 0.6705 - auc\_1: 0.7351 - val\_loss: 0.6924 - val\_accuracy: 0.6127 - val\_auc\_1: 0.6535

Epoch 5/5

476/476 [=====] - 5s 11ms/step - loss: 0.5868 - accuracy: 0.6817 - auc\_1: 0.7508 - val\_loss: 0.7127 - val\_accuracy: 0.6150 - val\_auc\_1: 0.6522

Out[211]:

<tensorflow.python.keras.callbacks.History at 0x7f28b2fe17f0>

## plot of the history

In [ ]:

```
# printing the plot of accuracy and loss and auc
it=model.history

import matplotlib.pyplot as plt
training_loss = it['loss']
test_loss = it['val_loss']
# Create count of the number of epochs
epoch_count = range(1, len(training_loss) + 1)
# Visualize loss history
plt.plot(epoch_count, training_loss, 'r--')
plt.plot(epoch_count, test_loss, 'b-')
plt.legend(['Training Loss', 'Test Loss'])
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title("PLOT OF LOSS ")
plt.show();
training_accuracy = it['auc']
test_accuracy = it['val_auc']
# Create count of the number of epochs
epoch_count = range(1, len(training_accuracy) + 1)
# Visualize loss history
plt.plot(epoch_count, training_accuracy, 'r--')
plt.plot(epoch_count, test_accuracy, 'b-')
plt.legend(['Training AUC', 'Test AUC'])
plt.xlabel('Epoch')
plt.ylabel('accuracy')
plt.title("PLOT OF AUC ")
plt.show()
training_accuracy1 = it['accuracy']
test_accuracy1 = it['val_accuracy']
# Create count of the number of epochs
epoch_count = range(1, len(training_accuracy1) + 1)
# Visualize loss history
plt.plot(epoch_count, training_accuracy1, 'r--')
plt.plot(epoch_count, test_accuracy1, 'b-')
plt.legend(['Training ACCURACY', 'Test ACCURACY'])
plt.xlabel('Epoch')
plt.ylabel('accuracy')
plt.title("PLOT OF ACCURACY")
plt.show();
```

